

boston-aiml

March 17, 2024

```
[1]: import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns
sns.set_theme(style = "darkgrid")

import warnings
warnings.filterwarnings("ignore", category = Warning)

from sklearn.datasets import load_boston
boston = load_boston()
```

#Make sure the data is loaded correctly.

```
[2]: data_url = "http://lib.stat.cmu.edu/datasets/boston"
col_names = boston.feature_names
df = pd.DataFrame(boston.data, columns = col_names)
df['MEDV'] = boston.target

# Make sure the data is loaded correctly.
df.head()
```

```
[2]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	

	PTRATIO	B	LSTAT	MEDV
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	5.33	36.2

#1.The model possible for this house price prediction are “Linear regression model”, “Random forest”.

#The output for this problem is expected to be a scalar (a constant).

1 2.Model Input / Output

```
[3]: x = df.iloc[:, :-1]
     y = df["MEDV"]

     # According to the defenition in this dataset, it is suitable to define "MEDV"
     ↪as output and others variables as input features.
```

2 3.EDA

```
[4]: print(df.describe())
fig, ax = plt.subplots(ncols = 7, nrows = 2, figsize = (12, 12))
index = 0
ax = ax.flatten()
for k, v in df.items():
    sns.boxplot(y = k, data = df, ax = ax[index])
    index += 1

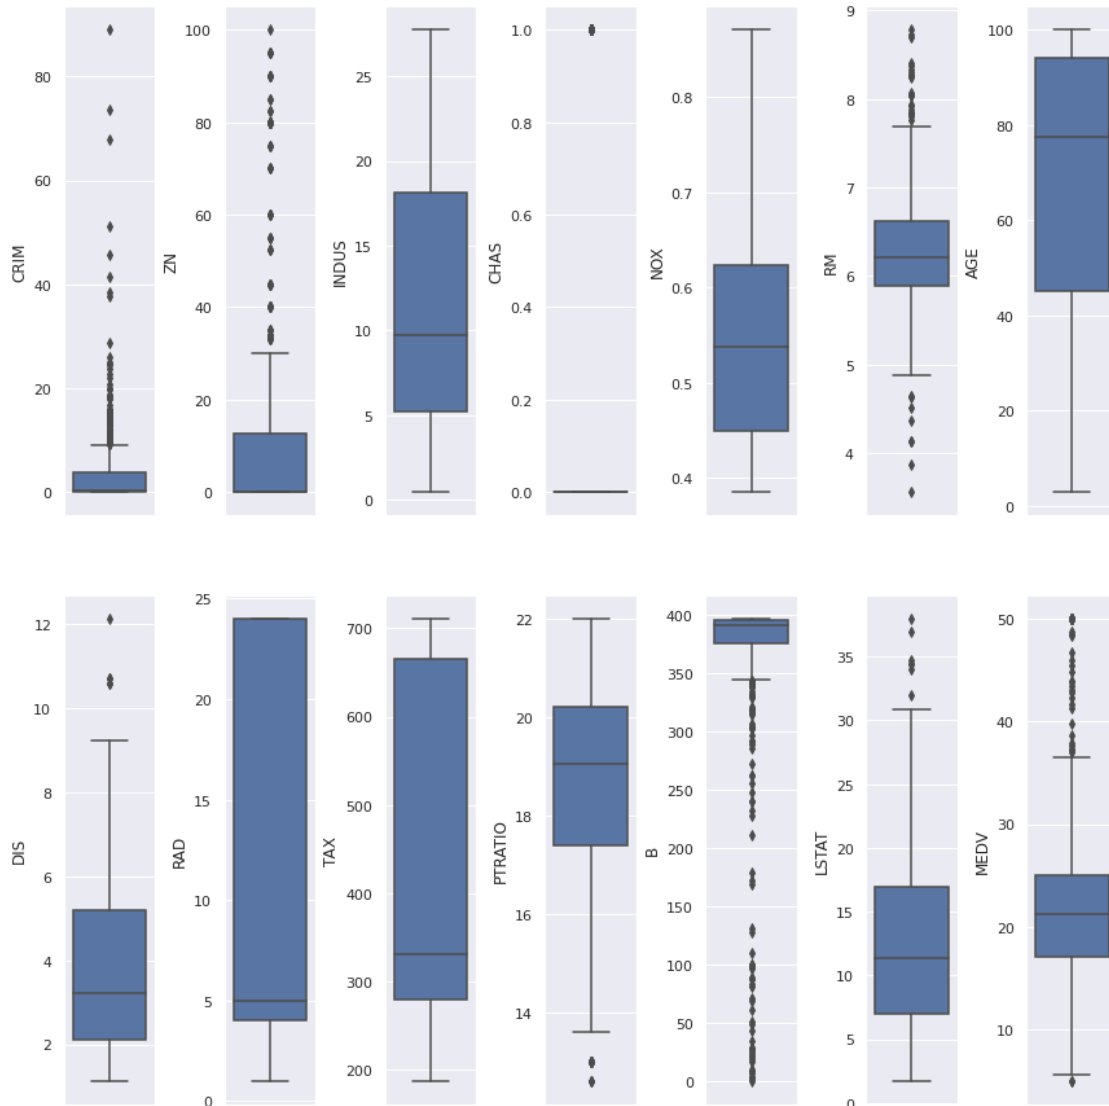
plt.tight_layout(pad = 0.4, w_pad = 0.5, h_pad = 5.0)

# According to the result of brief EDA, CHAS is a dummy, while the others are
↪numerical feartures.
# The variaties of the features are quite high; therefore, regularization is
↪required.
```

	CRIM	ZN	INDUS	CHAS	NOX	RM \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000

	AGE	DIS	RAD	TAX	PTRATIO	B \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	68.574901	3.795043	9.549407	408.237154	18.455534	356.674032
std	28.148861	2.105710	8.707259	168.537116	2.164946	91.294864
min	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000
25%	45.025000	2.100175	4.000000	279.000000	17.400000	375.377500
50%	77.500000	3.207450	5.000000	330.000000	19.050000	391.440000
75%	94.075000	5.188425	24.000000	666.000000	20.200000	396.225000
max	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000

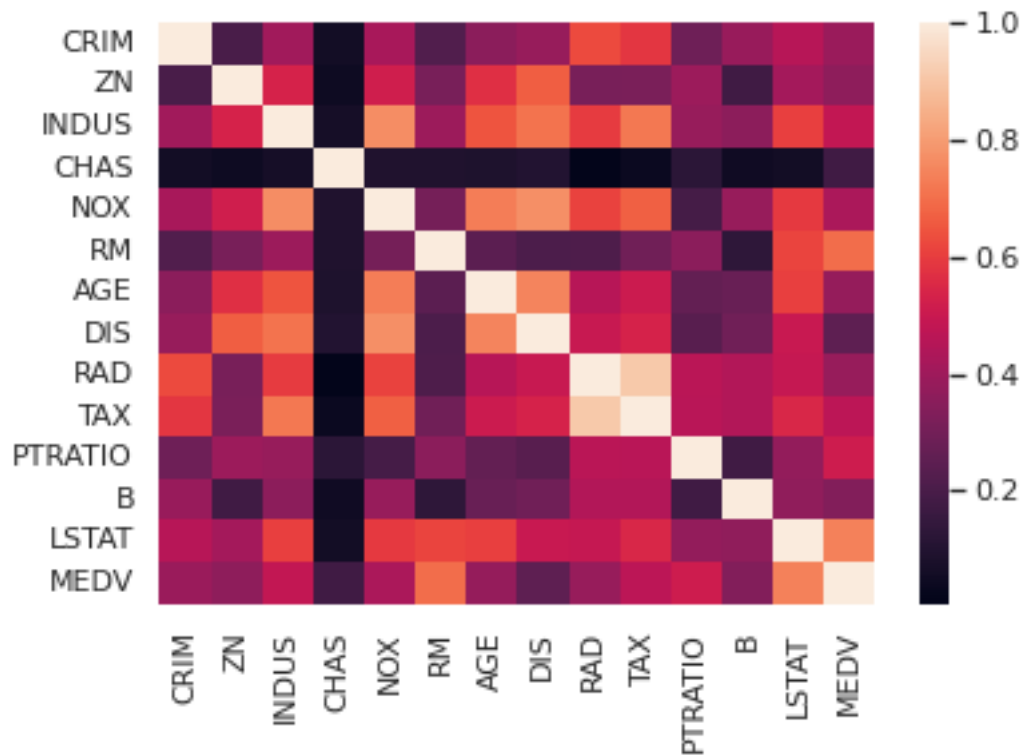
	LSTAT	MEDV
count	506.000000	506.000000
mean	12.653063	22.532806
std	7.141062	9.197104
min	1.730000	5.000000
25%	6.950000	17.025000
50%	11.360000	21.200000
75%	16.955000	25.000000
max	37.970000	50.000000



```
[5]: # Absolute correlation coefficients among all variables.
corr = df.corr().abs()
sns.heatmap(corr)
```

```
# It seems that (besides dummy) some of the features are correlated to
↳ themselves.
```

```
[5]: <AxesSubplot:>
```



```
[6]: # There's no missing value in the dataset.
df.isnull().sum()
```

```
[6]: CRIM      0
     ZN       0
     INDUS    0
     CHAS     0
     NOX      0
     RM       0
     AGE      0
     DIS      0
     RAD      0
     TAX      0
     PTRATIO  0
     B        0
     LSTAT    0
     MEDV     0
```

dtype: int64

```
[7]: # DataPreprocessing

# Ignore correlation coefficients which are too low to the output(set at 0.35).
cor_limit = 0.35
corr_list = x[x.columns[:]].apply(lambda x : x.corr(y))
cols = []
for i in corr_list.index:
    if (corr_list[i] >= cor_limit or corr_list[i] <= -cor_limit):
        cols.append(i)

# New data(x, input) with high correlation (HC) coefficients with Y.
HC_x = df[cols]
HC_x.head()
```

```
[7]:
```

	CRIM	ZN	INDUS	NOX	RM	AGE	RAD	TAX	PTRATIO	LSTAT
0	0.00632	18.0	2.31	0.538	6.575	65.2	1.0	296.0	15.3	4.98
1	0.02731	0.0	7.07	0.469	6.421	78.9	2.0	242.0	17.8	9.14
2	0.02729	0.0	7.07	0.469	7.185	61.1	2.0	242.0	17.8	4.03
3	0.03237	0.0	2.18	0.458	6.998	45.8	3.0	222.0	18.7	2.94
4	0.06905	0.0	2.18	0.458	7.147	54.2	3.0	222.0	18.7	5.33

```
[8]: # Regularization. For the large variaties in the input features.
# MaxAbsScaler - For the highest explain power among all scalers from view of PCA and high predition accuracy.

from sklearn.preprocessing import MaxAbsScaler
Z = HC_x

scaler = MaxAbsScaler()
scaler.fit(Z)
Z = scaler.transform(Z)

HC_x.iloc[:,:] = Z

"""
#decomposition(PCA if needed)
from sklearn.decomposition import PCA

pca = PCA(n_components = 2)
L = pca.fit_transform(Z)

#explain percentage
pca_explained = PCA()
pca_explained.fit(Z)
```

```
p_pca = np.round(pca_explained.explained_variance_ratio_, 2)
print("The 2 components explain percentage : " +
      str(round(p_pca[0]+p_pca[1],4)))
"""
```

```
[8]: '\n#decomposition(PCA if needed)\nfrom sklearn.decomposition import PCA\n\nnpca =
PCA(n_components = 2)\nL = pca.fit_transform(Z)\n\n\n#explain
percentage\nnpca_explained = PCA()\npca_explained.fit(Z)\np_pca =
np.round(pca_explained.explained_variance_ratio_, 2)\nprint("The 2 components
explain percentage : " + str(round(p_pca[0]+p_pca[1],4)))\n'
```

```
[9]: # Make sure the input is regularized.
HC_x.describe()
```

```
[9]:
```

	CRIM	ZN	INDUS	NOX	RM	AGE \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	0.040612	0.113636	0.401470	0.636849	0.715790	0.685749
std	0.096672	0.233225	0.247309	0.133040	0.080025	0.281489
min	0.000071	0.000000	0.016583	0.442021	0.405581	0.029000
25%	0.000922	0.000000	0.187094	0.515499	0.670330	0.450250
50%	0.002883	0.000000	0.349315	0.617681	0.707118	0.775000
75%	0.041327	0.125000	0.652487	0.716418	0.754385	0.940750
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

	RAD	TAX	PTRATIO	LSTAT
count	506.000000	506.000000	506.000000	506.000000
mean	0.397892	0.574173	0.838888	0.333238
std	0.362802	0.237042	0.098407	0.188071
min	0.041667	0.263010	0.572727	0.045562
25%	0.166667	0.392405	0.790909	0.183039
50%	0.208333	0.464135	0.865909	0.299184
75%	1.000000	0.936709	0.918182	0.446537
max	1.000000	1.000000	1.000000	1.000000

3 4.Training set and Testing set seperation. (by 15%)

```
[10]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(HC_x, y, test_size = 0.15)
```

4 5.Fitting of each ML model.

```
[11]: from sklearn import metrics
from sklearn.linear_model import LinearRegression

# Create a Linear Regression Model.
lm = LinearRegression()
```

```
lm.fit(x_train, y_train)
LinearRegression(fit_intercept=True, n_jobs=None)
```

[11]: LinearRegression()

```
[12]: # Model fitting on training data.
y_pred_lm = lm.predict(x_train)
print('R^2:', round(metrics.r2_score(y_train, y_pred_lm), 4))
print('Adjusted R^2:', round(1 - (1-metrics.r2_score(y_train, y_pred_lm)) * (len(y_train) - 1) / (len(y_train) - x_train.shape[1] - 1), 4))
print('MSE:', round(metrics.mean_squared_error(y_train, y_pred_lm), 4))
```

R^2: 0.6993

Adjusted R^2: 0.6921

MSE: 26.1001

```
[13]: # Predicting with the testing data.
y_test_pred_lm = lm.predict(x_test)

# Model Evaluation
score_linreg = metrics.r2_score(y_test, y_test_pred_lm)
print('R^2:', round(score_linreg, 4))
print('Adjusted R^2:', round(1 - (1-metrics.r2_score(y_test, y_test_pred_lm)) * (len(y_test) - 1) / (len(y_test) - x_test.shape[1] - 1), 4))
print('MSE:', round(metrics.mean_squared_error(y_test, y_test_pred_lm), 4))
```

R^2: 0.6793

Adjusted R^2: 0.63

MSE: 22.7771

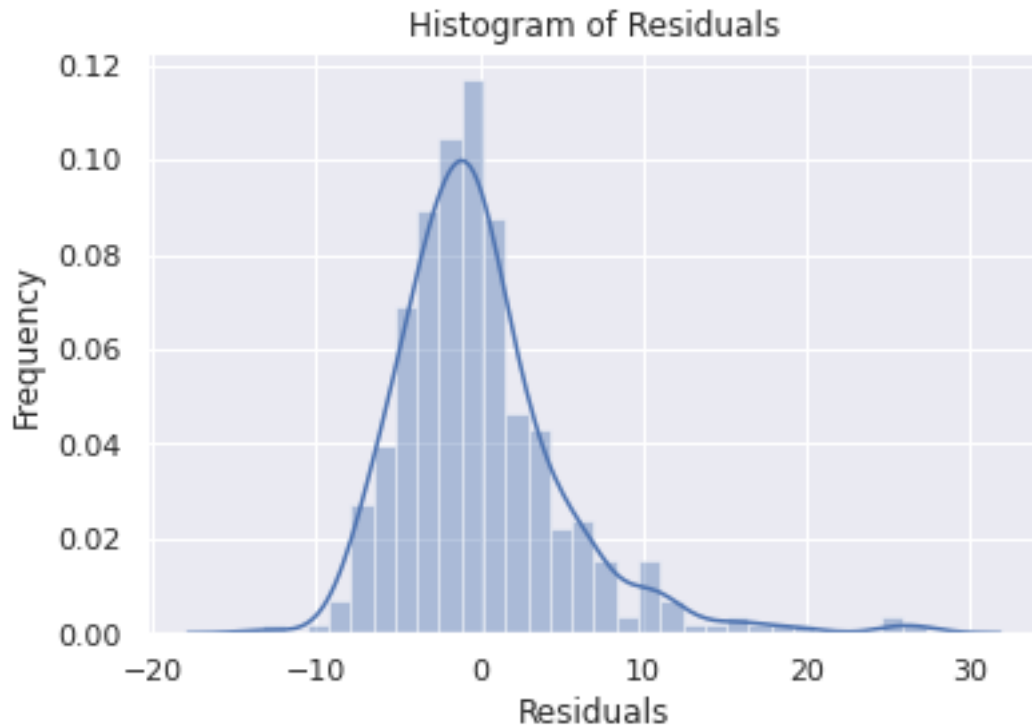
```
[14]: # Visualizing the differences between real prices and predicted prices.
plt.scatter(y_train, y_pred_lm)
plt.xlabel("Real Prices")
plt.ylabel("Predicted Prices")
plt.title("Real Prices vs Predicted prices")
plt.show()
```



```
[15]: # Checking of Normality of residuals.
from scipy import stats
statsn, pval = stats.normaltest(y_train - y_pred_lm)
print("p_val of normality test is : " + str(pval) + ", so it's normally_
distributed." + "\n")

sns.distplot(y_train - y_pred_lm)
plt.title("Histogram of Residuals")
plt.xlabel("Residuals")
plt.ylabel("Frequency")
plt.show()
```

p_val of normality test is : 8.049803151117528e-36, so it's normally distributed.



```
[16]: # Import Random Forest Regressor.
from sklearn.ensemble import RandomForestRegressor

# Create a Random Forest Regressor.
rfm = RandomForestRegressor()
rfm.fit(x_train, y_train)
```

```
[16]: RandomForestRegressor()
```

```
[17]: # Model fitting on training data.
y_pred_rf = rfm.predict(x_train)
print('R^2:',round(metrics.r2_score(y_train, y_pred_rf),4))
print('Adjusted R^2:',round(1 - (1-metrics.r2_score(y_train, y_
    ↪ y_pred_rf))*(len(y_train)-1)/(len(y_train)-x_train.shape[1]-1),4))
print('MSE:',round(metrics.mean_squared_error(y_train, y_pred_rf),4))
```

```
R^2: 0.976
Adjusted R^2: 0.9754
MSE: 2.083
```

```
[18]: # Predicting with the testing data.
y_test_pred_rf = rfm.predict(x_test)
```

```
# Model Evaluation
score_rf = round(metrics.r2_score(y_test, y_test_pred_rf),4)
print('R^2:', score_rf)
print('Adjusted R^2:',round(1 - (1-metrics.r2_score(y_test, y_test_pred_rf))*(len(y_test)-1)/(len(y_test)-x_test.shape[1]-1),4))
print('MSE:',round(metrics.mean_squared_error(y_test, y_test_pred_rf),4))
```

R^2: 0.8663

Adjusted R^2: 0.8458

MSE: 9.4939

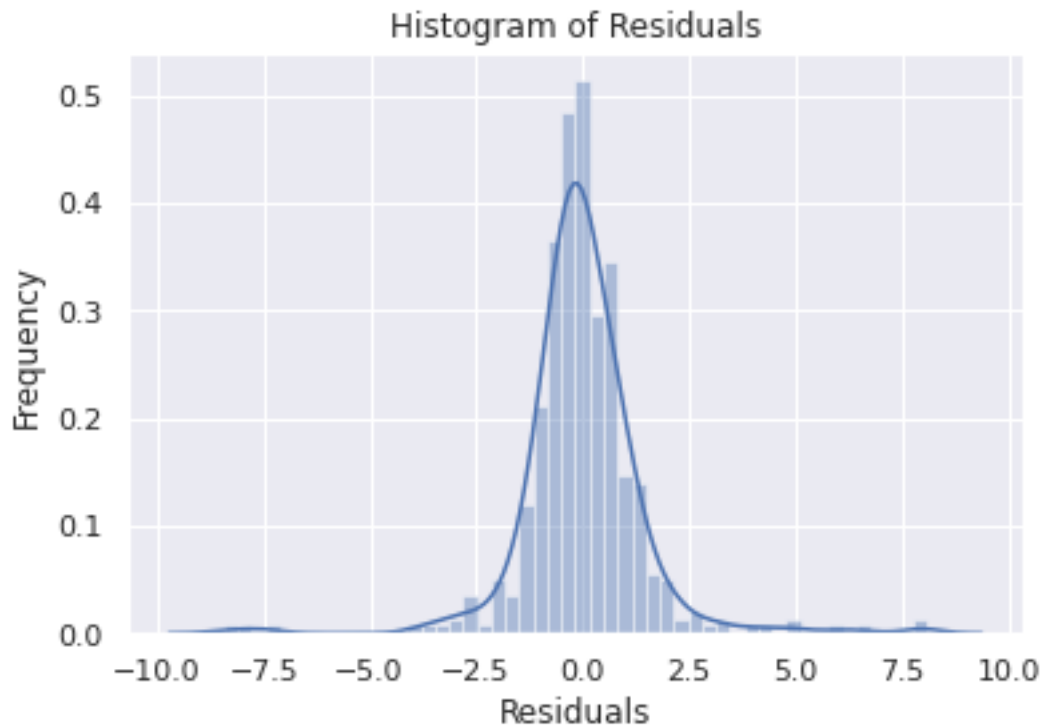
```
[19]: # Visualizing the differences between actual prices and predicted values
plt.scatter(y_train, y_pred_rf)
plt.xlabel("Real Prices")
plt.ylabel("Predicted Prices")
plt.title("Real Prices vs Predicted prices")
plt.show()
```



```
[20]: # Checking of Normality of residuals.
statsn, pval = stats.normaltest(y_train - y_pred_rf)
print("p_val of normality test is : " + str(pval) + ", so it's normally distributed." + "\n")
```

```
sns.distplot(y_train - y_pred_rf)
plt.title("Histogram of Residuals")
plt.xlabel("Residuals")
plt.ylabel("Frequency")
plt.show()
```

p_val of normality test is : 1.3311902577279215e-21, so it's normally distributed.



```
[21]: # Comparison of Linear Regression model and Random Forest Regressor.
# It seems that both of the model do not occur serious overfitting or underfitting problems.
# Random Forest Regressor seems to be outperforming the Linear Regression model.
```

5. Compare the accuracy with the regression models.

```
[22]: print("The definition of Acc. Here is referring to R-squared score.")
print("Acc. of Linear Regression Model : " + str(round(score_linreg, 4)))
print("Acc. of Random Forest Model : " + str(round(score_rf, 4)))
```

The definition of Acc. Here is referring to R-squared score.
Acc. of Linear Regression Model : 0.6793

Acc. of Random Forest Model : 0.8663

6 6.Neural Nnetwork

```
[23]: import keras

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import Dropout
from sklearn.metrics import r2_score

# Since the prediction accuracy for the output of a scalar would be 0 for most
↳ of the scenarios.
# Therefore, I decided to define the accuracy of NN model as  $R^2$  as the
↳ classical linear regression model.
"""
from tensorflow.keras import backend as K

def ssr(y_true, y_pred):
    ss_res = K.sum(K.square(y_true - y_pred))
    ss_tot = K.sum(K.square(y_true - K.mean(y_true)))
    return(1 - (ss_res / (ss_tot + K.epsilon()))))
"""
```

```
[23]: '\nfrom tensorflow.keras import backend as K\n\ndef ssr(y_true, y_pred):\n\n    ss_res = K.sum(K.square(y_true - y_pred))\n    ss_tot = K.sum(K.square(y_true -\n    K.mean(y_true)))\n    return(1 - (ss_res / (ss_tot + K.epsilon())))\n'
```

```
[24]: # NN Model stacking.
model_nn = Sequential()
model_nn.add(Dense(10, input_dim = len(HC_x.columns)))
model_nn.add(Dense(10))
model_nn.add(Dense(10))
model_nn.add(Dense(1))
model_nn.compile(loss = "mse", optimizer = "rmsprop", metrics = [r2_score],
↳ run_eagerly = True)

model_nn.summary()
```

2022-12-05 13:17:17.832387: I
tensorflow/core/common_runtime/process_util.cc:146] Creating new thread pool
with default inter op setting: 2. Tune using inter_op_parallelism_threads for
best performance.

Model: "sequential"

Layer (type)	Output Shape	Param #
--------------	--------------	---------

```

=====
dense (Dense)                (None, 10)                110
-----
dense_1 (Dense)              (None, 10)                110
-----
dense_2 (Dense)              (None, 10)                110
-----
dense_3 (Dense)              (None, 1)                 11
=====
Total params: 341
Trainable params: 341
Non-trainable params: 0
-----

```

```

[25]: # NN Model training.
      # This dataset is a small one, which may somehow occur overfitting problem.

nn_result = model_nn.fit(x_train, y_train, batch_size = 20, epochs = 30,
      ↪validation_split = 0.10)

fig, ax1 = plt.subplots()
ax2 = ax1.twinx()

ax1.plot(nn_result.history["r2_score"], color = "g", label = "R2")
ax2.plot(nn_result.history["loss"], color = "b", label = "Loss")

ax1.set_ylabel("r2_score", color = "g")
ax2.set_ylabel("loss", color = "b")

plt.title("Model Accuracy")
plt.xlabel("epoch")
fig.legend()
plt.show()

print("Brief view on the result of training.")

```

```

2022-12-05 13:17:18.279095: I
tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:185] None of the MLIR
Optimization Passes are enabled (registered 2)

```

```

Epoch 1/30
20/20 [=====] - 5s 60ms/step - loss: 550.2806 -
r2_score: -6.7891 - val_loss: 509.3455 - val_r2_score: -4.6462
Epoch 2/30
20/20 [=====] - 1s 57ms/step - loss: 468.0719 -
r2_score: -5.7202 - val_loss: 425.5466 - val_r2_score: -3.4199
Epoch 3/30
20/20 [=====] - 1s 57ms/step - loss: 373.9273 -

```

```

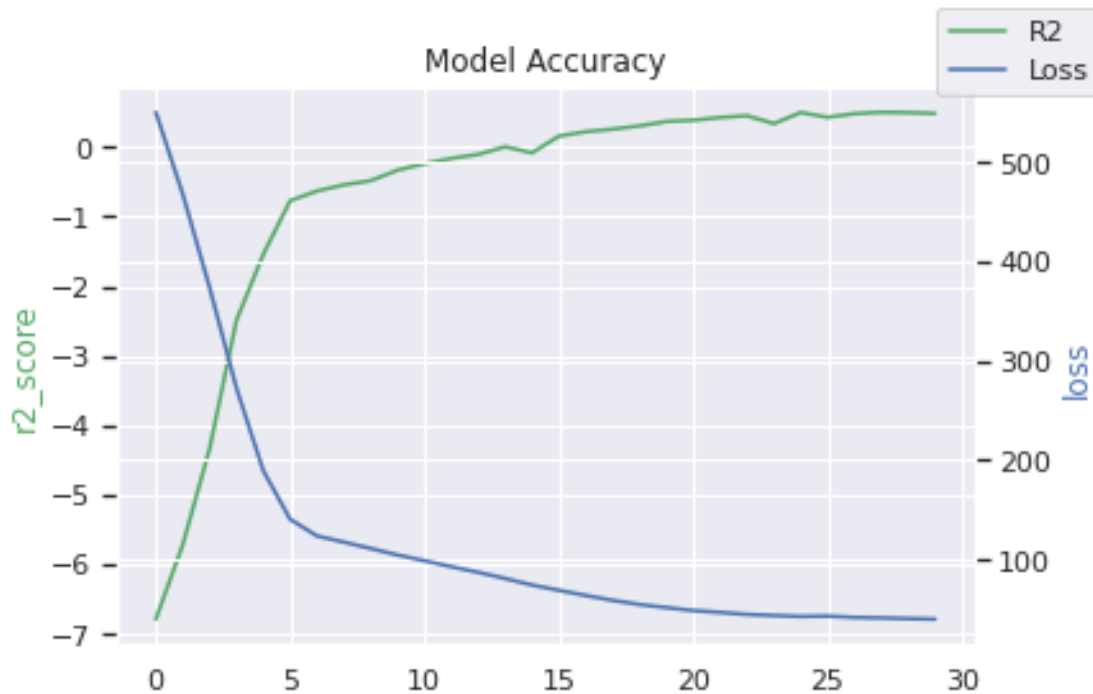
r2_score: -4.3381 - val_loss: 325.2667 - val_r2_score: -2.0890
Epoch 4/30
20/20 [=====] - 1s 57ms/step - loss: 272.6080 -
r2_score: -2.4913 - val_loss: 232.3489 - val_r2_score: -1.1219
Epoch 5/30
20/20 [=====] - 1s 57ms/step - loss: 189.3460 -
r2_score: -1.5415 - val_loss: 169.6804 - val_r2_score: -0.8865
Epoch 6/30
20/20 [=====] - 1s 56ms/step - loss: 140.4627 -
r2_score: -0.7720 - val_loss: 141.1949 - val_r2_score: -1.3927
Epoch 7/30
20/20 [=====] - 1s 56ms/step - loss: 123.6893 -
r2_score: -0.6321 - val_loss: 134.1029 - val_r2_score: -1.2913
Epoch 8/30
20/20 [=====] - 1s 57ms/step - loss: 117.5221 -
r2_score: -0.5445 - val_loss: 127.1136 - val_r2_score: -1.3974
Epoch 9/30
20/20 [=====] - 1s 59ms/step - loss: 111.0973 -
r2_score: -0.4840 - val_loss: 120.5698 - val_r2_score: -1.1983
Epoch 10/30
20/20 [=====] - 1s 57ms/step - loss: 104.5461 -
r2_score: -0.3325 - val_loss: 114.5736 - val_r2_score: -1.3373
Epoch 11/30
20/20 [=====] - 1s 56ms/step - loss: 98.7414 -
r2_score: -0.2400 - val_loss: 107.4133 - val_r2_score: -0.9971
Epoch 12/30
20/20 [=====] - 1s 57ms/step - loss: 92.7498 -
r2_score: -0.1641 - val_loss: 101.6103 - val_r2_score: -1.0256
Epoch 13/30
20/20 [=====] - 1s 57ms/step - loss: 87.0843 -
r2_score: -0.1068 - val_loss: 94.0513 - val_r2_score: -0.6189
Epoch 14/30
20/20 [=====] - 1s 57ms/step - loss: 80.8990 -
r2_score: 0.0031 - val_loss: 87.0259 - val_r2_score: -0.4735
Epoch 15/30
20/20 [=====] - 1s 57ms/step - loss: 74.3350 -
r2_score: -0.0844 - val_loss: 80.4614 - val_r2_score: -0.1584
Epoch 16/30
20/20 [=====] - 1s 56ms/step - loss: 69.0157 -
r2_score: 0.1561 - val_loss: 75.0943 - val_r2_score: 0.0083
Epoch 17/30
20/20 [=====] - 1s 56ms/step - loss: 63.7904 -
r2_score: 0.2196 - val_loss: 67.5616 - val_r2_score: -0.1237
Epoch 18/30
20/20 [=====] - 1s 57ms/step - loss: 58.8932 -
r2_score: 0.2561 - val_loss: 61.2925 - val_r2_score: 0.0610
Epoch 19/30
20/20 [=====] - 1s 59ms/step - loss: 54.7316 -

```

```

r2_score: 0.3030 - val_loss: 56.9293 - val_r2_score: 0.1464
Epoch 20/30
20/20 [=====] - 1s 57ms/step - loss: 51.6036 -
r2_score: 0.3669 - val_loss: 52.6521 - val_r2_score: 0.2800
Epoch 21/30
20/20 [=====] - 1s 57ms/step - loss: 48.6013 -
r2_score: 0.3839 - val_loss: 50.2821 - val_r2_score: 0.5113
Epoch 22/30
20/20 [=====] - 1s 57ms/step - loss: 46.7983 -
r2_score: 0.4243 - val_loss: 47.6687 - val_r2_score: 0.5754
Epoch 23/30
20/20 [=====] - 1s 57ms/step - loss: 44.6932 -
r2_score: 0.4489 - val_loss: 43.6461 - val_r2_score: 0.5680
Epoch 24/30
20/20 [=====] - 1s 57ms/step - loss: 43.6698 -
r2_score: 0.3358 - val_loss: 42.9058 - val_r2_score: 0.6188
Epoch 25/30
20/20 [=====] - 1s 56ms/step - loss: 42.6587 -
r2_score: 0.4956 - val_loss: 42.0499 - val_r2_score: 0.5124
Epoch 26/30
20/20 [=====] - 1s 57ms/step - loss: 43.0076 -
r2_score: 0.4265 - val_loss: 40.7416 - val_r2_score: 0.6342
Epoch 27/30
20/20 [=====] - 1s 61ms/step - loss: 41.7197 -
r2_score: 0.4792 - val_loss: 40.5412 - val_r2_score: 0.6510
Epoch 28/30
20/20 [=====] - 1s 60ms/step - loss: 41.2645 -
r2_score: 0.4988 - val_loss: 39.6165 - val_r2_score: 0.6539
Epoch 29/30
20/20 [=====] - 1s 57ms/step - loss: 40.5412 -
r2_score: 0.4940 - val_loss: 39.2298 - val_r2_score: 0.6590
Epoch 30/30
20/20 [=====] - 1s 57ms/step - loss: 39.9451 -
r2_score: 0.4814 - val_loss: 37.5831 - val_r2_score: 0.6471

```



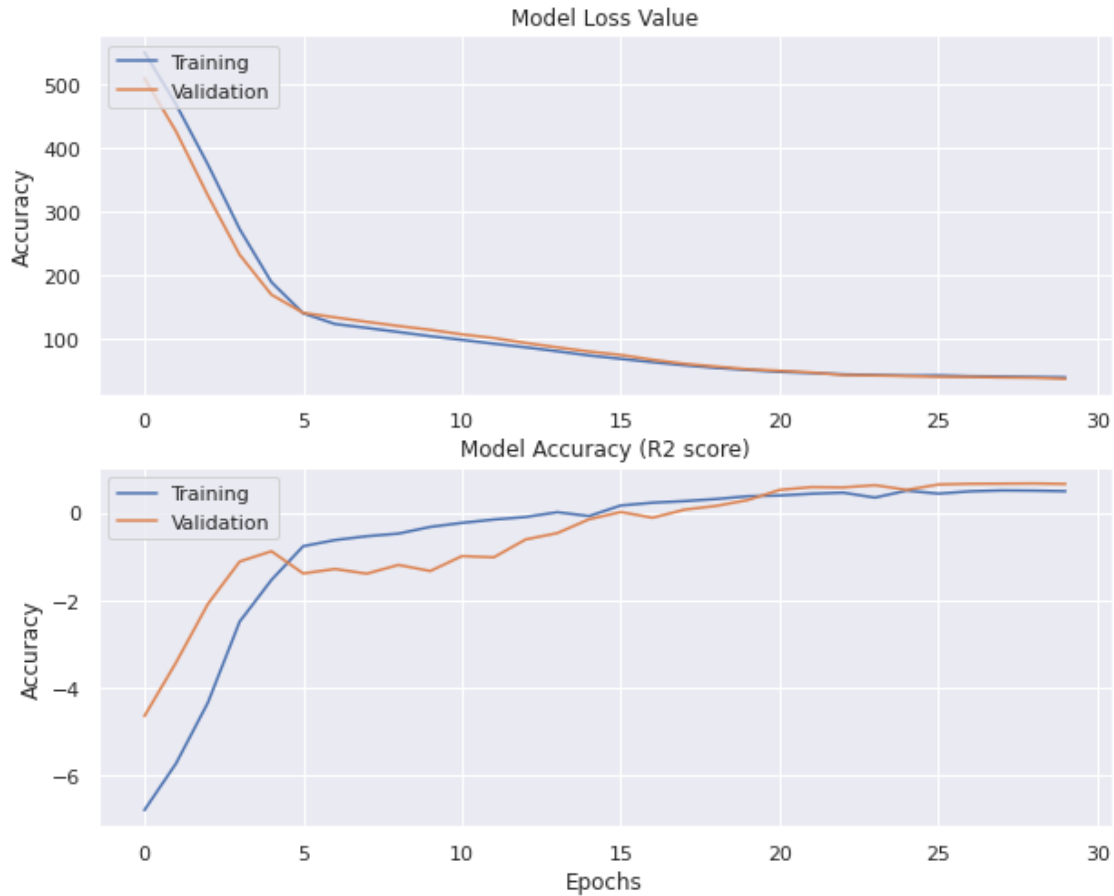
Brief view on the result of training.

```
[26]: # Examine the process of NN_model on loss and R2 score.
plt.figure(figsize = (10, 8))
plt.subplot(211)

plt.plot(nn_result.history["loss"])
plt.plot(nn_result.history["val_loss"])
plt.title("Model Loss Value")
plt.ylabel("Accuracy")
plt.legend(["Training", "Validation"], loc = "upper left")

plt.subplot(212)

plt.plot(nn_result.history["r2_score"])
plt.plot(nn_result.history["val_r2_score"])
plt.title("Model Accuracy (R2 score)")
plt.ylabel("Accuracy")
plt.xlabel("Epochs")
plt.legend(["Training", "Validation"], loc = "upper left")
plt.show()
```

```
[27]: # Accuracy of well-trainde NN Model.
scores = model_nn.evaluate(x_test, y_test)
print("Accuracy of CNN is : " + str(round(scores[1],4)))
#The accurcyt is quite unstable...
```

3/3 [=====] - 0s 38ms/step - loss: 38.5716 - r2_score: 0.3765

Accuracy of CNN is : 0.3765

```
[28]: # Compare the accuracy of all the regression models.
print("The defenition of Acc. here is referring to R-squared score.")
print("Acc. of Linear Regression Model : " + str(round(score_linreg, 4)))
print("Acc. of Random Forest Model : " + str(round(score_rf, 4)))
print("Acc. of Neural Network Model : " + str(round(scores[1],4)))
```

The defenition of Acc. here is referring to R-squared score.

Acc. of Linear Regression Model : 0.6793

Acc. of Random Forest Model : 0.8663

Acc. of Neural Network Model : 0.3765